

# How to use the standard Helm Charts

## Overview

To standardize how microservices are deployed into Kubernetes, the RDE Deploy team manages a repository of Helm Charts for WBD applications. This document includes a FAQ for how to best use the microservice Helm chart.

## FAQ

### Helm Chart Logistics Questions

Where do I find documentation for the Helm Chart?

- Answer...

The [README](#) in the helm-chart repository contains a table showing all the configurable variables the chart uses, along with the default value, and description/example. Clicking on the hyperlink in the `key` will take you to where that value is in the values file. Be sure to scroll the window to the right to view the description/example

How can I change the Helm Chart version for my microservice?

- Answer...

As part of the RDE release and deploy workflow, your microservice is deployed with a helm chart that has some version associated with it. This version can be specified for your entire service and can be specified for individual components. To do this, you must have the `helm-chart` and `helm-chart-version` make targets defined in your Makefile which are responsible for outputting a string representing the name and semver of the helm chart to use respectively. See [the example set up](#) in `wbd-golang-template` repository. All components will use the values defined in the root Makefile if they are not defined in its Makefile. **We highly recommend you using the latest available Helm chart version by setting the version value to `^0` (example).** This is the default setup from `wbd-golang/java-template`.

How can I render a Helm Chart locally?

- Answer...

Follow the guide [here](#) on pulling the `microservice` helm chart locally from Artifactory and testing it against your values files.

### Common Configuration Questions

How can I utilize an AWS role to access AWS infrastructure?

- [Helm Chart Deployment Failure due to Inability to Access AWS Resources](#)

How can I specify my own readiness, liveness, or startup probe block?

- Answer...

You can use the `.probe.[liveness|readiness|startup].override` to specify a full Kubernetes probe block, the documentation for which you can find here [Configure Liveness, Readiness and Startup Probes](#). For example to use a custom liveness probe but the standard readiness probe you can specify something like the following in your values file:

```
1 probe:
2   liveness:
3     override:
4       httpGet:
5         path: /livez
6         port: 9090
```

```
7     httpHeaders:
8         - name: Custom-Header
9             value: Awesome
10        initialDelaySeconds: 3
11        periodSeconds: 3
12    readiness:
13        path: /healthz
14        initialDelaySeconds: 2
```

This will render out:

```
1 ---
2 # Source: microservice/templates/deployment.yaml
3 ....
4     containers:
5         ...
6         livenessProbe:
7             httpGet:
8                 httpHeaders:
9                     - name: Custom-Header
10                        value: Awesome
11                        path: /healthz
12                        port: 9090
13                initialDelaySeconds: 3
14                periodSeconds: 3
15         readinessProbe:
16             httpGet:
17                 path: /healthz
18                 port: 8080
19             initialDelaySeconds: 2
20 ...
```

How do I set up Ingress?

▼ Answer...

Refer here: <https://wbdstreaming.atlassian.net/wiki/spaces/PEO/pages/89522269/BOLT+Onboarding+Handbook#10.-Boundary-Services-details>

How do I enable gRPC?

▼ Answer...

Set the following value.

```
1 enableGrpc: true
```

Also see: [How do I change the container port and service ports?](#)

How do I change the container port and service ports?

▼ Answer...

To change container and service port configurations, your values file must contain the following:

```
1 containerPort: # Container Port configuration
2     http: 8080
3     grpc: 18080
4
5 service:      # Service port configuration
```

```
6   type: ClusterIP
7   ports:
8     http: 80
9     grpc: 180
```

How do I enable HPA and set CPU/memory threshold?

▼ Answer...

HPA ([see documentation](#)) can be configured with the following values:

```
1 autoscaling:
2   enabled: true
3   minReplicas: 1
4   maxReplicas: 100
5   targetCPUUtilizationPercentage: 80
6   targetMemoryUtilizationPercentage: 80
```

Both `targetCPUUtilizationPercentage` and `targetMemoryUtilizationPercentage` are optional and will create an element of `type: Resource` under `HorizontalPodAutoscaler.spec.metrics`.

**\*IMPORTANT\***: A `resource request` **MUST** be defined for the HPA to work. The utilization percentage is calculated based on the amount of cpu/memory requested. You can set a resource block with the following values:

```
1 # example
2 resources:
3   requests:
4     memory: "64Mi"
5     cpu: "250m"
6   limits:
7     memory: "128Mi"
8     cpu: "500m"
```

**NOTE:** `targetCPUUtilizationPercentage` is always set in the `autoscaling` value. To unset it, set `targetCPUUtilizationPercentage: null`

How do I configure HPA scaleUp/scaleDown behavior?

▼ Answer...

**\*VERSION 0.16.1 AND LATER\***

A full `HPA behavior` block can be defined under `autoscaling`. Example:

```
1 autoscaling:
2   enabled: true
3   minReplicas: 1
4   maxReplicas: 100
5   targetCPUUtilizationPercentage: 80
6   behavior:
7     scaleDown:
8       policies:
9         - type: Percent
10           value: 10
11           periodSeconds: 60
12     scaleUp:
13       stabilizationWindowSeconds: 2
14       policies:
15         - type: Percent
16           value: 100
17           periodSeconds: 15
```

How do I configure HPA with custom metrics?

- Answer...

SRE team has started rolling out KEDA to support this capability. To learn more, click here: [KEDA - HELM Configuration](#)

#### \*NO LONGER SUPPORTED\*

This was previously implemented by a tool that's no longer supported by newer versions of Kubernetes. The SRE Team is currently reviewing a tool called KEDA to bring back custom autoscaling capabilities

#### \*OLD DOCUMENTATION\*

The chart supports HPA scaling on other metrics besides CPU and memory. See this [walkthrough on how to autoscale on multiple metrics and custom metrics](#). Users can create metrics for type `Pod`, `External`, and `Object` using the values `podMetrics`, `externalMetrics`, and `objectMetrics`. These values expect a list of a block for `metrics.(pods|external|object).*`. See [API docs](#) for what different metric types require.

Example values.yaml	Rendered hpa.yaml by Helm
<pre> 1  autoscaling: 2    enabled: true 3    minReplicas: 2 4    maxReplicas: 30 5    targetCPUUtilizationPercentage: null 6    podMetrics:          # Two different metrics 7      - metric: 8        name: packets_per_second 9        target: 10       type: AverageValue 11       averageValue: 1k 12      - metric: 13        name: some_other_pod_metric 14        target: 15       type: AverageUtilization 16       averageUtilization: 50 17 18    externalMetrics:          # One 'External' metric 19      - metric: 20        name: queue_messages_ready 21        selector: 22          matchLabels: 23            queue: "worker_tasks" 24        target: 25          type: AverageValue 26          averageValue: 30 27 28    objectMetrics:          # One 'Object' metric 29      - metric: 30        name: requests-per-second 31        describedObject: 32          apiVersion: networking.k8s.io/v1 33          kind: Ingress </pre>	<pre> 1  apiVersion: autoscaling/v2 2  kind: HorizontalPodAutoscaler 3  metadata: 4    ... 5  spec: 6    ... 7    minReplicas: 2 8    maxReplicas: 30 9    metrics: 10   - type: Pods 11     pods: 12       metric: 13         name: packets_per_second 14         target: 15           averageValue: 1k 16           type: AverageValue 17   - type: Pods 18     pods: 19       metric: 20         name: some_other_pod_metric 21         target: 22           averageUtilization: 50 23           type: AverageUtilization 24   - type: External 25     external: 26       metric: 27         name: queue_messages_ready 28         selector: 29           matchLabels: 30             queue: worker_tasks 31         target: 32           averageValue: 30 33           type: AverageValue </pre>

```

34     name: main-route
35   target:
36     type: Value
37     value: 10k

```

```

34   - type: Object
35     object:
36       describedObject:
37         apiVersion: networking.k8s.io/v1
38         kind: Ingress
39         name: main-route
40       metric:
41         name: requests-per-second
42       target:
43         type: Value
44         value: 10k

```

#### How do I create a Job and CronJob?

▼ Answer...

Both CronJobs and Jobs are set with the same way. A list of batch jobs are defined in the `cronJobs` value. If the job contains a schedule, it runs as a cronJob, else, it runs as a job. The values that can be set are listed below:

Value	Required	Description	Default
<code>name</code>	Yes	Set to <code>.metadata.name</code>	None
<code>schedule</code>	No	CronJob Schedule. If omitted, creates a Job	None
<code>env</code>	No	Set environment variables in the pod	None
<code>envFieldRefs</code>	No	Set <a href="#">pod information</a> as environment variables	None
<code>args</code>	No	Arguments to run within a container	None
<code>successfulJobsHistoryLimit</code>	No	Sets <code>.spec.successfulJobsHistoryLimit</code>	5
<code>failedJobsHistoryLimit</code>	No	Sets <code>.spec.failedJobsHistoryLimit</code>	5
<code>suspend</code>	No	Sets <code>.spec.suspend</code>	<code>false</code>
<code>restartPolicy</code>	No	Sets <code>.spec.restartPolicy</code> within the Job template	<code>OnFailure</code>
<code>jobAnnotations</code>	No	Sets <code>.metadata.annotations</code> within the Job template	None

Below is an example on how to create a cronJob “my-cronjob” and Job “my-job”

```

1 # Toggle if you wish to turn off Deployments and Services
2 deploymentEnable: false
3 serviceEnable: false
4
5 cronJobs:

```

```

6 - name: "my-cronjob"                      # CronJob Creation
7   schedule: "* * * * *"
8   env:
9     NAME: val
10    envFieldRefs:
11      MY_POD_NAMESPACE:
12        apiVersion: v1
13        fieldPath: metadata.namespace
14    successfulJobsHistoryLimit: 10
15    failedJobsHistoryLimit: 10
16    suspend: true
17    restartPolicy: OnFailure
18    jobAnnotations:
19      annotation1: value
20 - name: "my-job"                          # Job Creation
21   env:
22     NAME_1: jobVal
23   args:
24     - argVal1
25     - argVal2
26   restartPolicy: Never
27   jobAnnotations:
28     annotation1: val1

```

How do I expose pod information to my microservice?

▼ Answer...

You can set pod information as environment variables. See documentation: [Downward API](#). To use this in the microservice Helm chart, use the `envFieldRefs` value.

Example:

```

1 envFieldRefs:
2   HOSTNAME:
3     apiVersion: v1
4     fieldPath: status.podIP
5   MY_POD_NAMESPACE:
6     apiVersion: v1
7     fieldPath: metadata.namespace

```

This sets the the `HOSTNAME` variable as the pod's IP address in and the `MY_POD_NAMESPACE` variable as the namespace. Including the `apiVersion` is optional. It defaults to `v1`

How can I set a JSON or YAML as an environment variable?

▼ Answer...

The chart supports a `toJson` and `toYaml` feature for environment variables.

Example 1:

```

1 env:
2   REGION: "us-east-1"
3   toYaml:
4     APP_YAML_CONFIG:
5       one: 1
6       two: 2
7     ANOTHER_APP_YAML_CONFIG:
8       five: 5
9       six: 6

```

```
10  toJson:
11    SPRING_APPLICATION_JSON:
12      {"seven": "7"}
13    ANOTHER_SPRING_APPLICATION_JSON:
14      {"eight": "8"}
```

How do I change the number of pods running for my microservice?

▼ Answer...

Set `replicaCount` to the desired number

Example:

```
1  replicaCount: 10
```

Note: this doesn't take effect when `autoscaling` is enabled.

How can I inject a file into a configMap?

▼ Answer...

Configuration files must be embedded within helm values files. As part of the [build-microservice-image-and-push](#) workflow, you can specify a path in your repository, and that file will be injected in into the values files prior to being uploaded to the values files being uploaded to Artifactory.

Values files can specify a `fileData` field in the configMap like so:

```
1  configMap:
2    enable: true
3    mountPath: /mount/path
4    data:
5      example_data: example_val
6    fileData:
7      - key: conf.json
8        path: ".../config.json"
9
10
```

The workflow will update the values file to be:

```
1  configMap:
2    enable: true
3    mountPath: /mount/path
4    data:
5      conf.json: |-
6        {
7          "data": "value"
8        }
9      example_data: example_val
10     fileData:
11       - key: conf.json
12         path: ".../config.json"
13
```

The `fileData` field isn't used at all by the helm chart itself.